

AFL for Linux Network Servers

Fuzzing  
For  
Worms.



Dobin Rutishauser



01

SSL/TLS Recommendations

// OWASP Switzerland 2013

02

BurpSentinel - Semi Automated Web Scanner

// BSides Vienna 2014

03

Automated WAF Testing and XSS Detection

// OWASP Switzerland 2015, Barcamp

04

Lecturer Exploit Development & Mitigations

// Berner Fachhochschule, 2016, 2017, 2018

**What we have**

01

File Based Fuzzer

**What I want**

02

Network Fuzzer

**Demo**

03

Praise the Demo gods!

**Feedback Driven Fuzzing**

04

Honggfuzz

**Results**

05

0days?

```
cd ~/binutils-2.25
```

```
CC=afl-gcc ./configure
```

```
make
```

```
mkdir afl_in afl_out
```

```
cp /bin/ps afl_in/
```

```
afl-fuzz -i afl_in -o afl_out  
./binutils/readelf -a @@
```

american fuzzy lop 1.86b (test)

process timing		overall results
run time	: 0 days, 0 hrs, 0 min, 2 sec	cycles done : 0
last new path	: none seen yet	total paths : 1
last uniq crash	: 0 days, 0 hrs, 0 min, 2 sec	uniq crashes : 1
last uniq hang	: none seen yet	uniq hangs : 0
cycle progress		map coverage
now processing	: 0 (0.00%)	map density : 2 (0.00%)
paths timed out	: 0 (0.00%)	count coverage : 1.00 bits/tuple
stage progress		findings in depth
now trying	: havoc	favorable paths : 1 (100.00%)
stage execs	: 1464/5000 (29.28%)	new edges on : 1 (100.00%)
total execs	: 1697	total crashes : 39 (1 unique)
exec speed	: 626.5/sec	total hangs : 0 (0 unique)
fuzzing strategy yields		path geometry
bit flips	: 0/16, 1/15, 0/13	levels : 1
byte flips	: 0/2, 0/1, 0/0	pending : 1
arithmetics	: 0/112, 0/25, 0/0	pend fav : 1
known ints	: 0/10, 0/28, 0/0	own finds : 0
dictionary	: 0/0, 0/0, 0/0	imported : n/a
havoc	: 0/0, 0/0	variable : 0
trim	: n/a, 0.00%	

[cpu: 92%]

IJG jpeg <a href="#">1</a>	libjpeg-turbo <a href="#">1</a> <a href="#">2</a>	libpng <a href="#">1</a>
libtiff <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a> <a href="#">5</a>	mozjpeg <a href="#">1</a>	PHP <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a> <a href="#">5</a> <a href="#">6</a> <a href="#">7</a> <a href="#">8</a>
Mozilla Firefox <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a>	Internet Explorer <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a>	Apple Safari <a href="#">1</a>
Adobe Flash / PCRE <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a> <a href="#">5</a> <a href="#">6</a> <a href="#">7</a>	sqlite <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a> ...	OpenSSL <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a> <a href="#">5</a> <a href="#">6</a> <a href="#">7</a>
LibreOffice <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a>	poppler <a href="#">1</a> <a href="#">2</a> ...	freetype <a href="#">1</a> <a href="#">2</a>
GnuTLS <a href="#">1</a>	GnuPG <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a>	OpenSSH <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a> <a href="#">5</a>
PuTTY <a href="#">1</a> <a href="#">2</a>	ntpd <a href="#">1</a> <a href="#">2</a>	nginx <a href="#">1</a> <a href="#">2</a> <a href="#">3</a>
bash (post-Shellshock) <a href="#">1</a> <a href="#">2</a>	tcpdump <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a> <a href="#">5</a> <a href="#">6</a> <a href="#">7</a> <a href="#">8</a> <a href="#">9</a>	JavaScriptCore <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a>
pdfium <a href="#">1</a> <a href="#">2</a>	ffmpeg <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a> <a href="#">5</a>	libmatroska <a href="#">1</a>
libarchive <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a> <a href="#">5</a> <a href="#">6</a> ...	wireshark <a href="#">1</a> <a href="#">2</a> <a href="#">3</a>	ImageMagick <a href="#">1</a> <a href="#">2</a> <a href="#">3</a> <a href="#">4</a> <a href="#">5</a> <a href="#">6</a> <a href="#">7</a> <a href="#">8</a> <a href="#">9</a> ...



Everything already  
fuzzed :-)

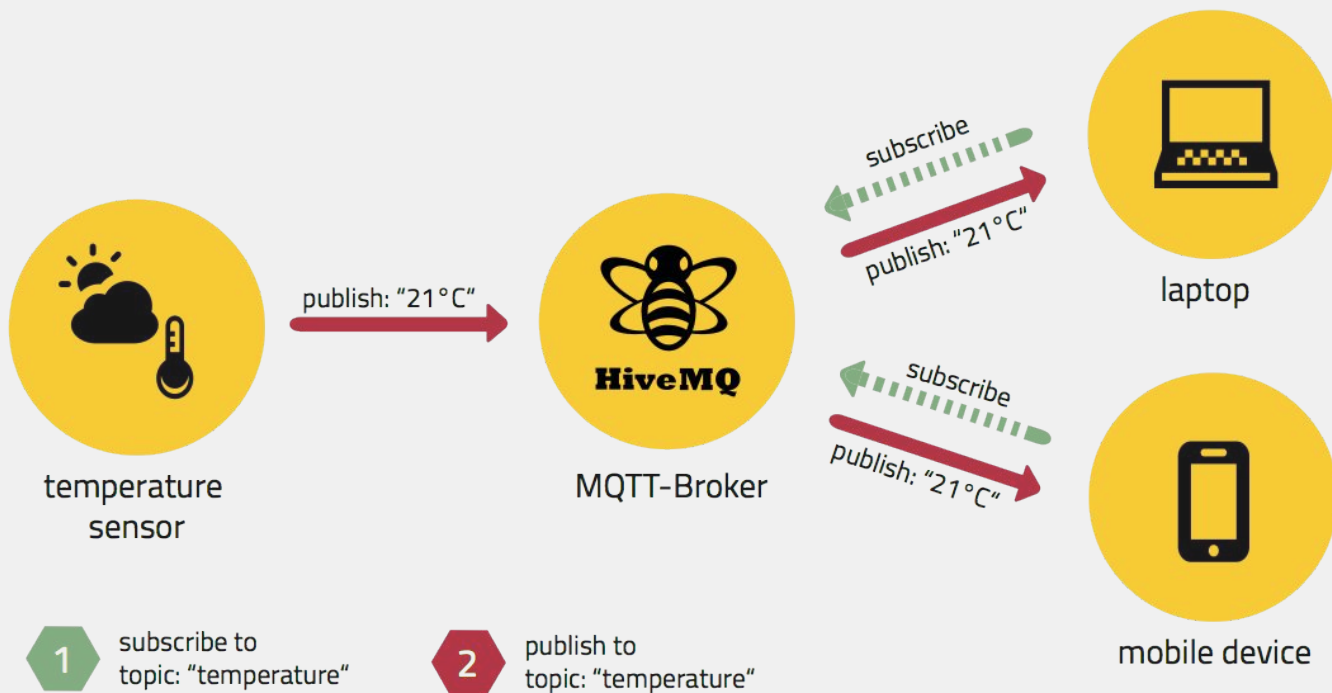


Why no fuzzing network  
servers?



OSI Model		
Layer		Protocol data unit (PDU)
Host layers	7. Application	Data
	6. Presentation	
	5. Session	
	4. Transport	Segment (TCP) / Datagram (UDP)
Media layers	3. Network	Packet
	2. Data link	Frame
	1. Physical	Bit





01 **CONNECT**

02 **SUBSCRIBE <channel>**

03 **PUBLISH <channel> <msg>**

# Using AFL to fuzz a MQTT server?



## Method 1: Wrap

Wrap packet-parsing  
function



## Method 2: Rewire

Re-wire read(), write()  
to read from files



## Method 1: Wrap

Isolate main functionality

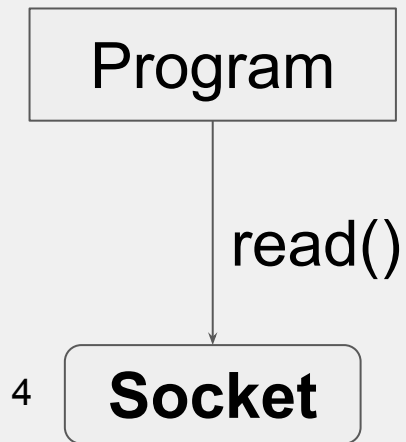
Write a wrapper

```
// Libfuzzer
```

```
// AFL Persistent Mode
```

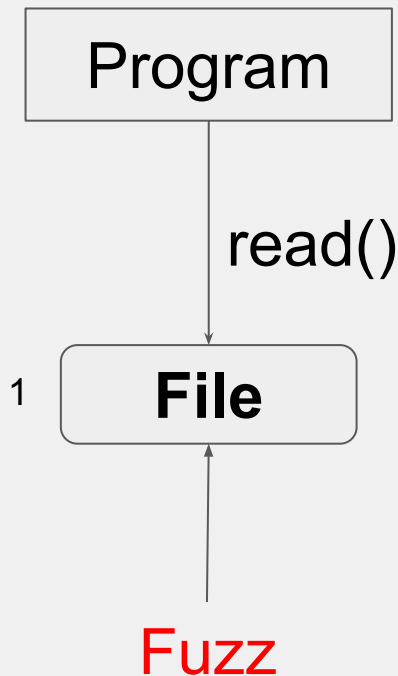
```
// WinAFL
```

```
extern "C" int LLVMFuzzerTestOneInput(  
    const uint8_t *data,  
    size_t size)  
{  
    parse_mqtt_packet(data);  
}
```



## Method 2: Re-wire

Re-wire `read()`, `write()`  
to read from files (not  
from sockets)



## Method 2: Re-wire

Re-wire `read()`, `write()`  
to read from files (not  
from sockets)

// AFL Style

// Preeny desock

# Case Study: **Mongoose 6.8**

CSNC-2017-023 //

Manual Re-Wire for de-socking

Fuzz with AFL

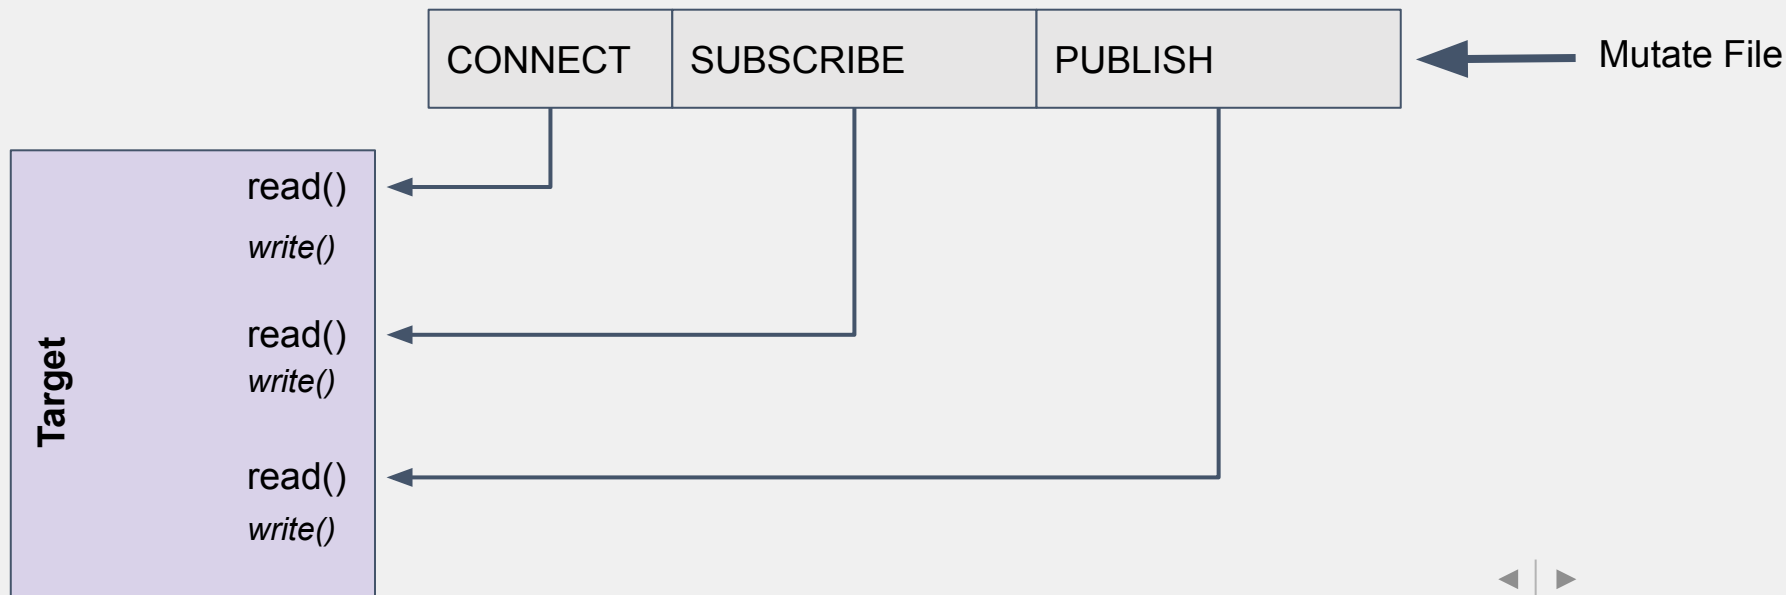
Stack BOF RCE Exploit





## File Based Fuzzer Problems

What are packets? //

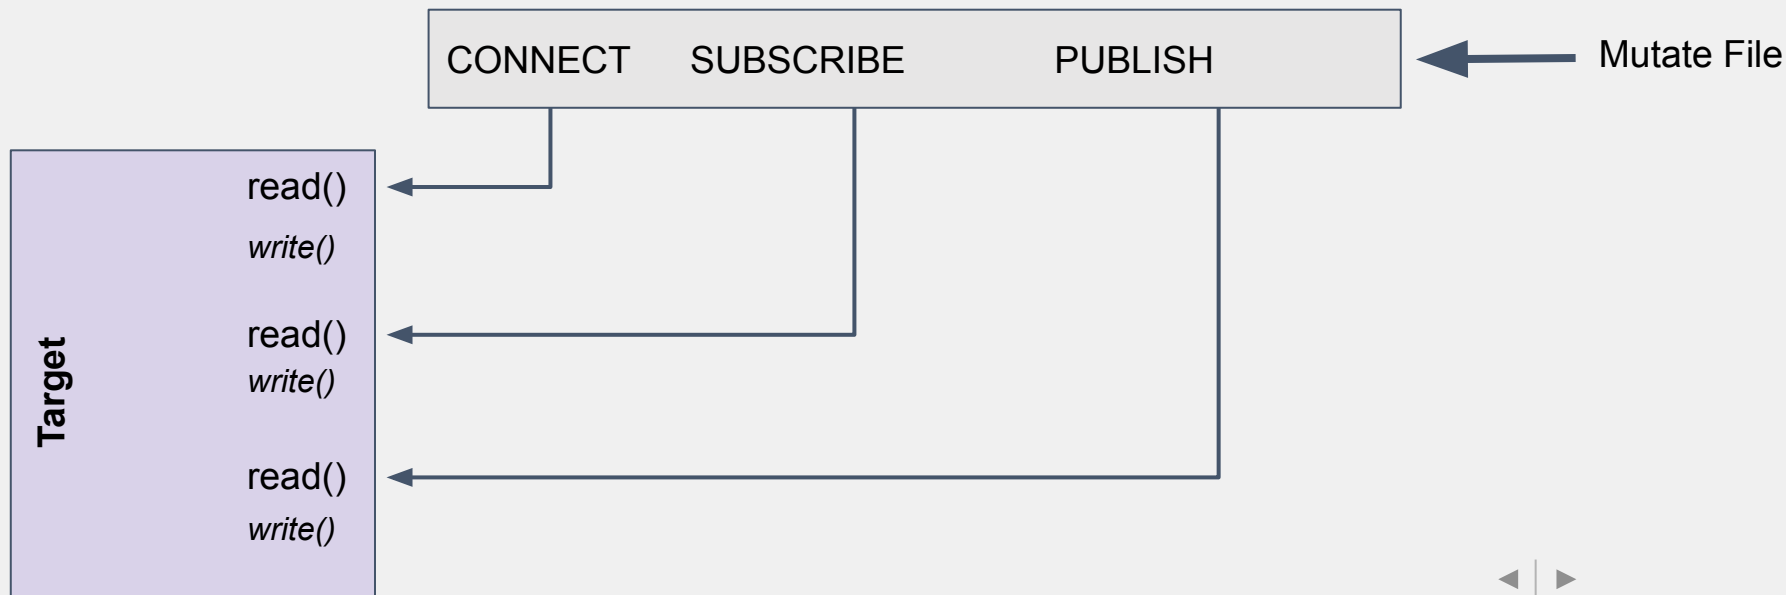






## File Based Fuzzer Problems

What are packets? //



# A better way? Generational Fuzzers

## Protocol Specifications //

Reproduce protocol in code as XML  
Specify every field, datatype, etc.

```
<DataModel name="DataPASV">
  <String value="PASV "/>
  <String value=""/>
  <String value="\r\n"/>
</DataModel>

<!-- TYPE [type_first] [type_second]:
set the type of file to be transferred -->
<DataModel name="DataTYPE1">
  <String value="TYPE "/>
  <String length="1" value="A"/>
  <Block minOccurs="0" maxOccurs="1">
    <String length="1" mutable="true">
      <Hint name="ValidValues" value="N;T;C"
    </String>
  </Block>
  <String value="\r\n"/>
</DataModel>
```

## Generators:

- Peachfuzz (XML “Pits”)
- Spike -> Sulley -> Boofuzz (python specs)
- Blab
- Dharma

## Protocol Learning:

- Netzob (UI)
- Pulsar (AI)



## Stateless

DNS, DHCP, HTTP, COAP

Ignore reply

Packets are independant

*Use AFL*

## Stateful

MQTT, VNC, Teamspeak

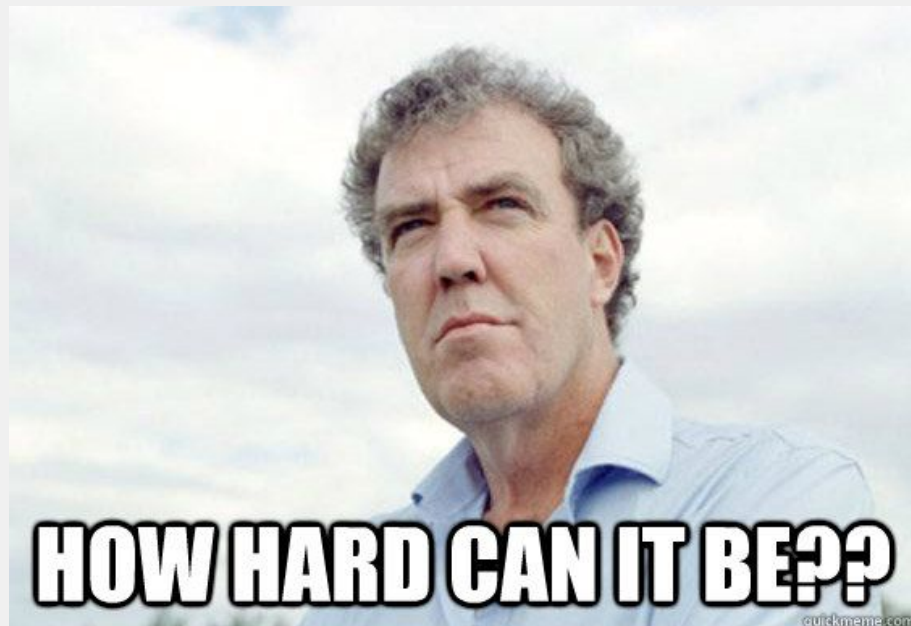
Responses important

Packets dependant

*Use Peach*

I dont want to **patch**  
I dont want to **XML** protocols

**i want to fuzz!**





01

## Intercept

Capture  
Network Data



02

## Replay Data

Fuzz some data  
and send



03

## Detect Crash

Check if server is  
still alive



04

## ???



05

## Profit!

FFW Basic  
Mongoose 6.8  
Demo

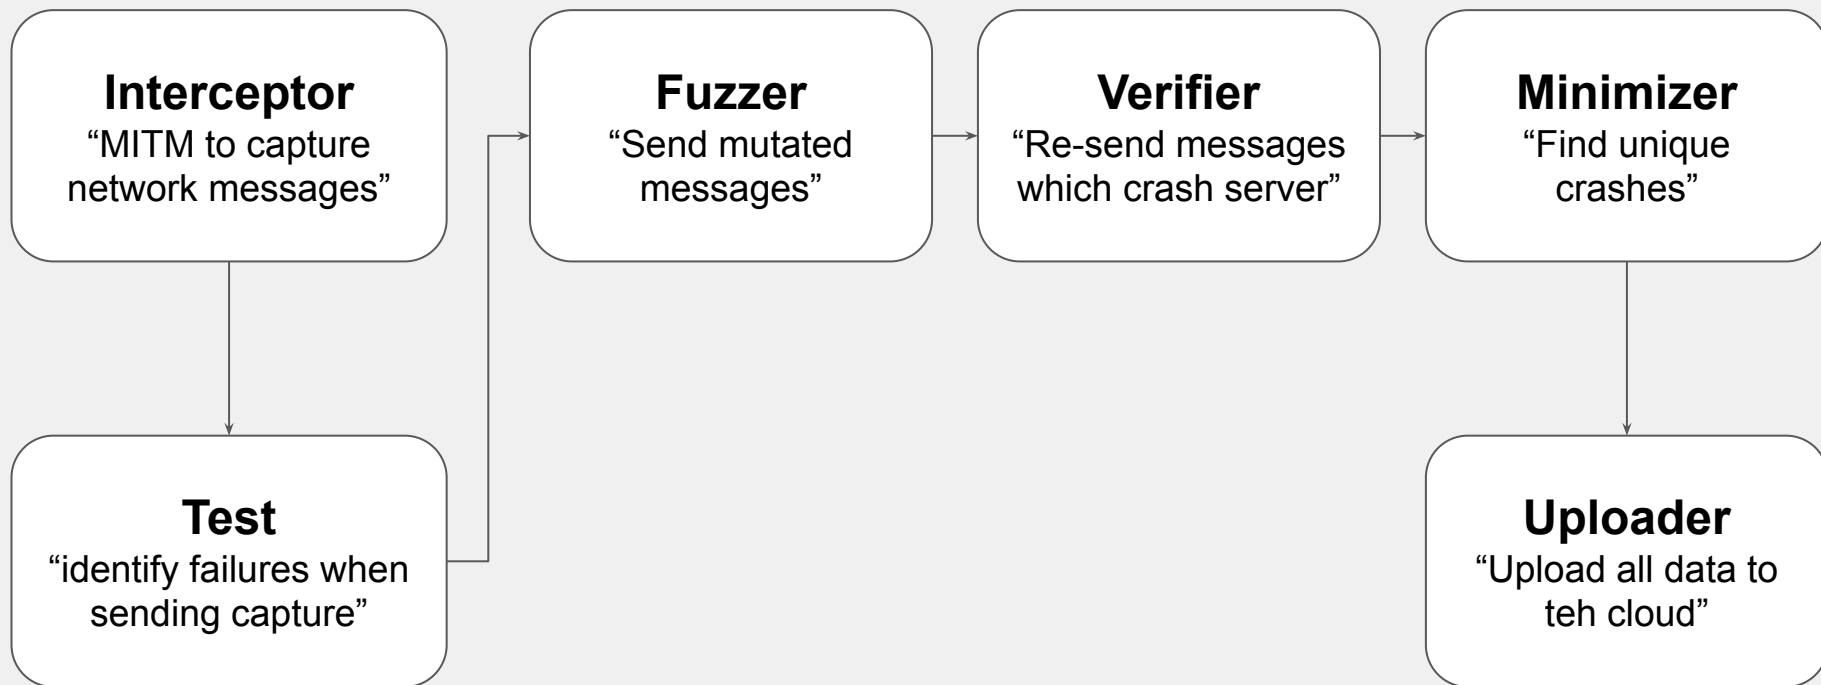
Praise  
the  
Demo  
Gods

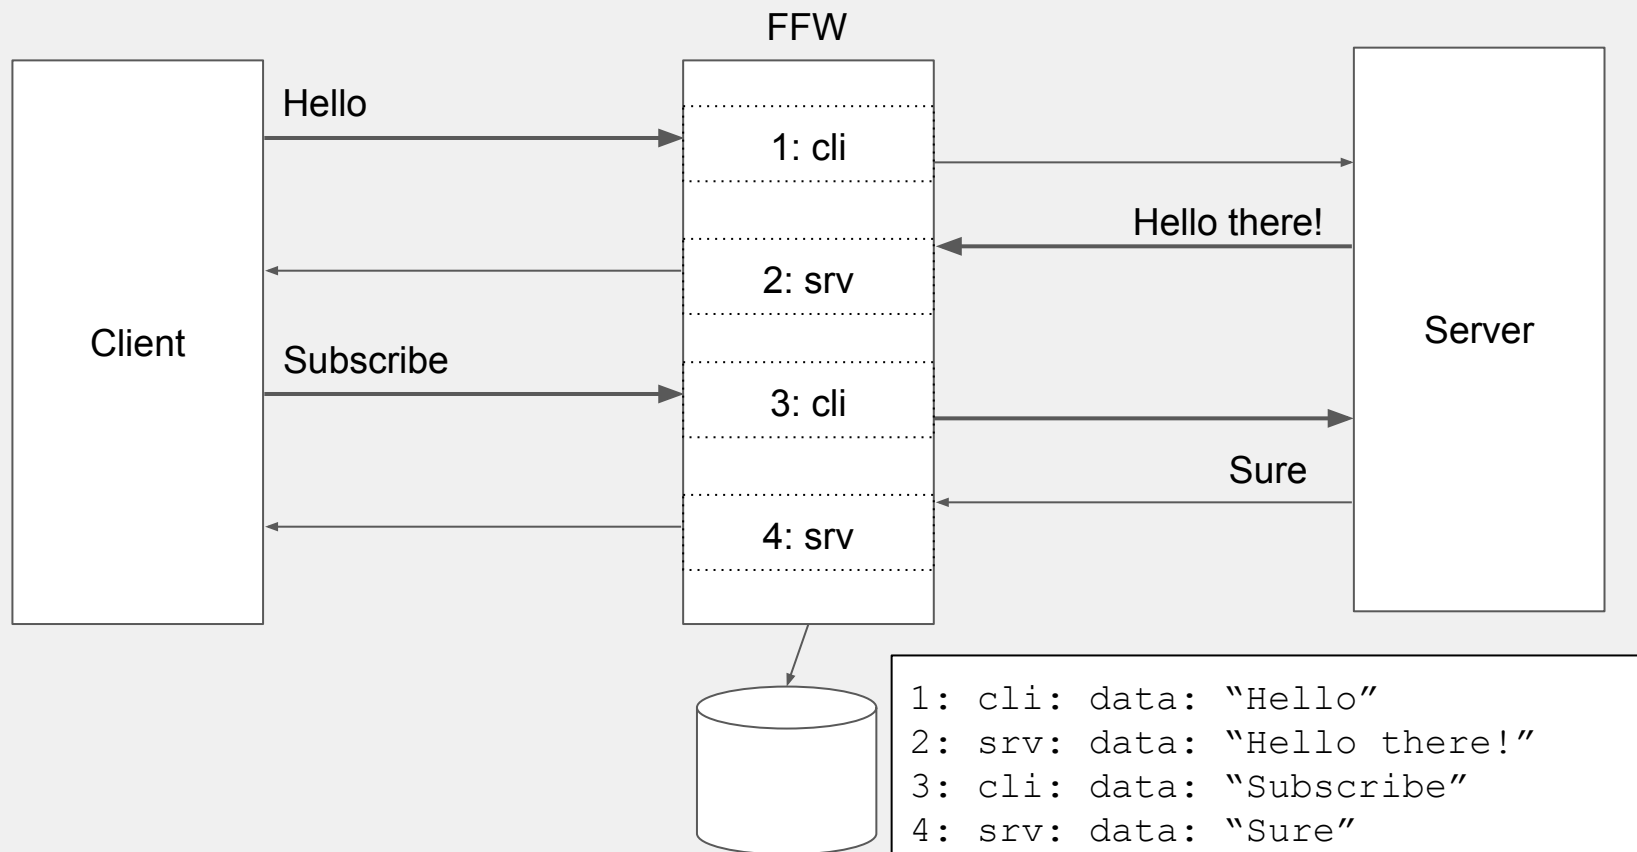


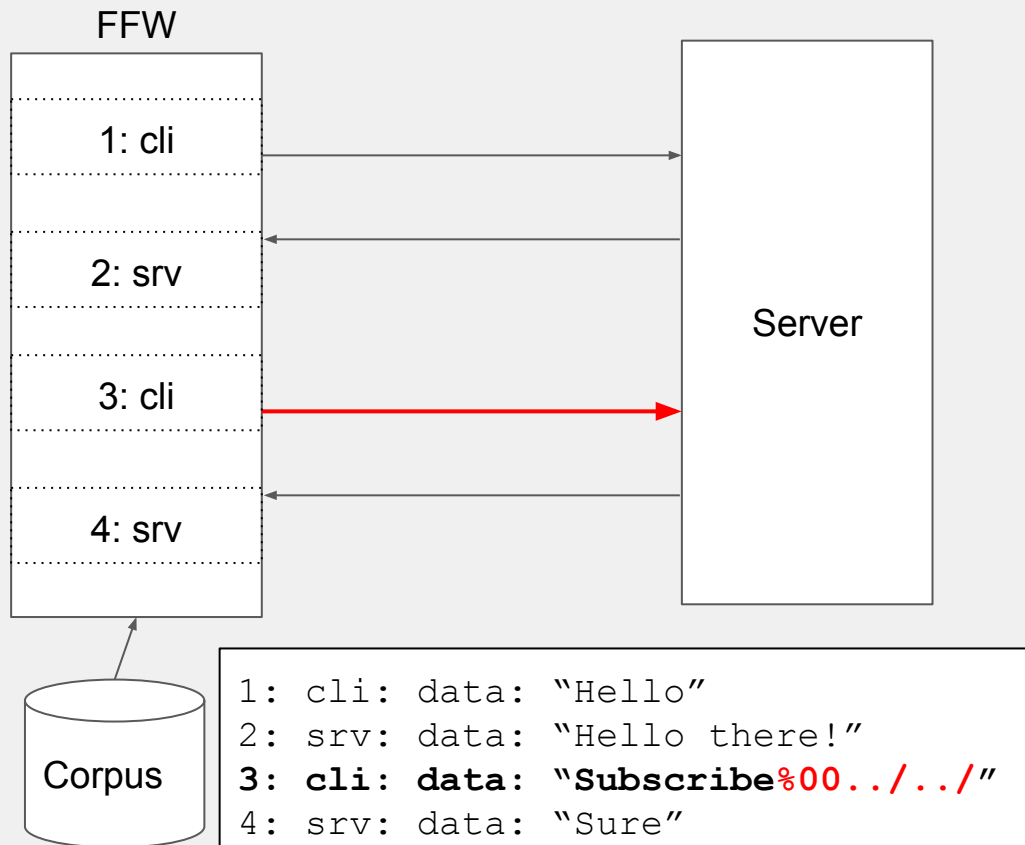
# FFW Internals







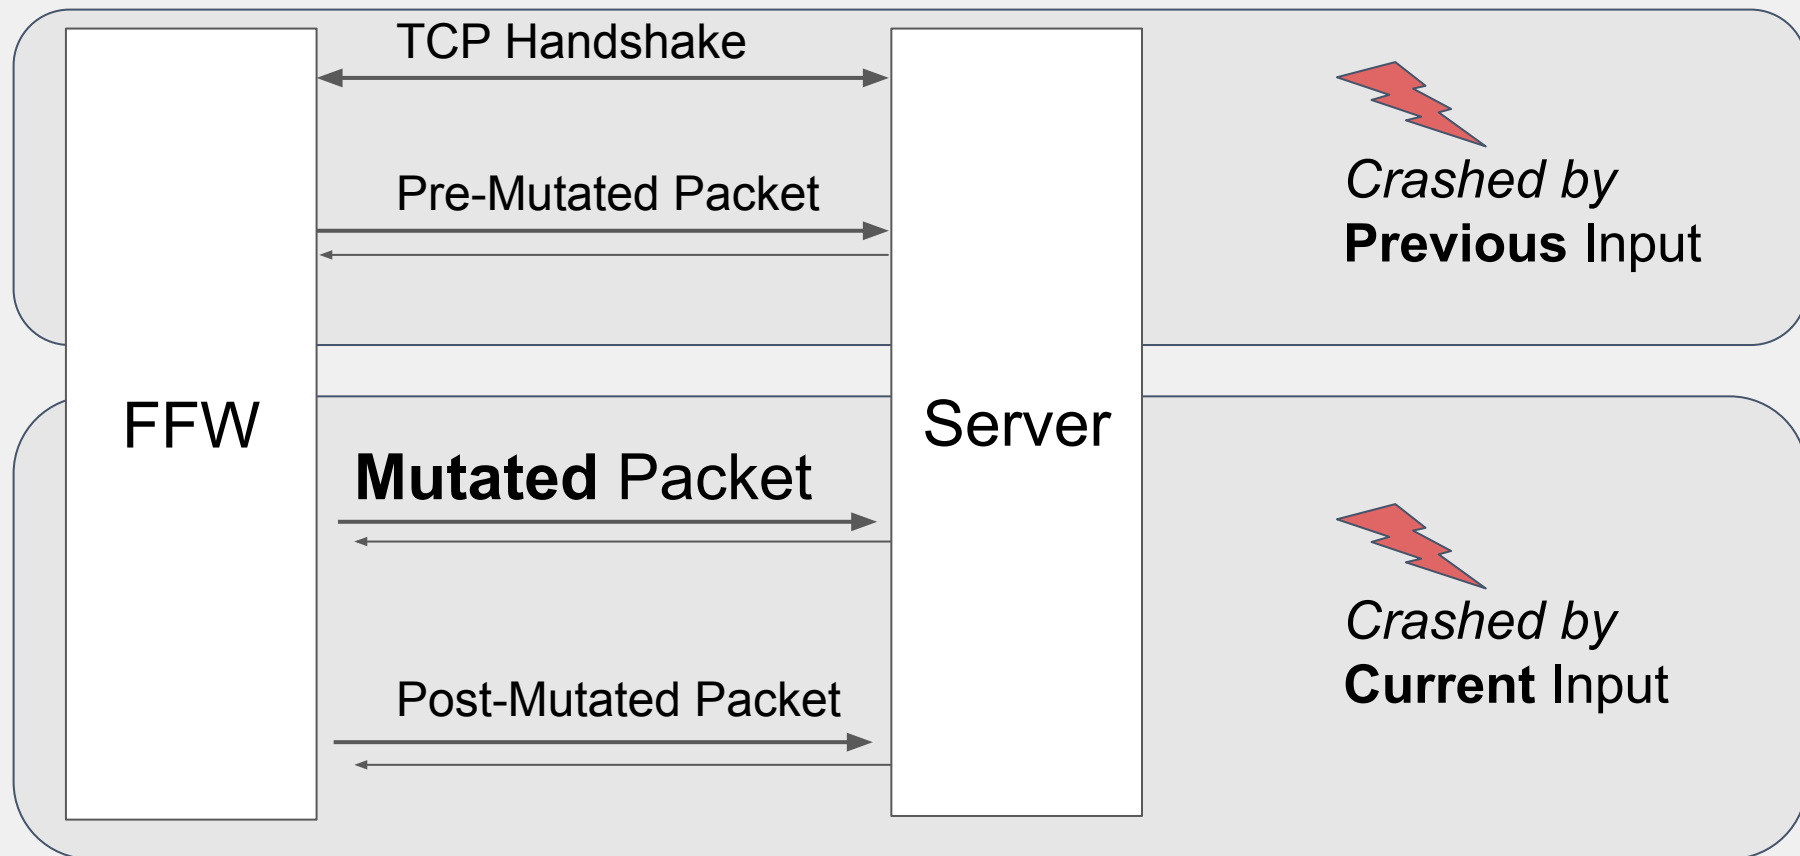




```
$ echo "aaa" | radamsa  
aaaa
```

Here radamsa decided to add one 'a' to the input. Let's try that again.

```
$ echo "aaa" | radamsa  
:aaa
```



## Crash Details: **Stack Trace**

01

PTRACE  
// **Bad, Unreliable**

02

ASAN  
// **Non-Stack BoF only**

03

GDB  
// **Need Parser...**

# Fuzzing Hardware

## What I have //

**Xeon E5 2670** (from 2012)

20mb Cache!

**8 cores, 16 threads!**

32 GB RAM!

2.6ghz!

**x2! rawr!**

777\$, [natex.us](http://natex.us)



## Prozessor



992.-

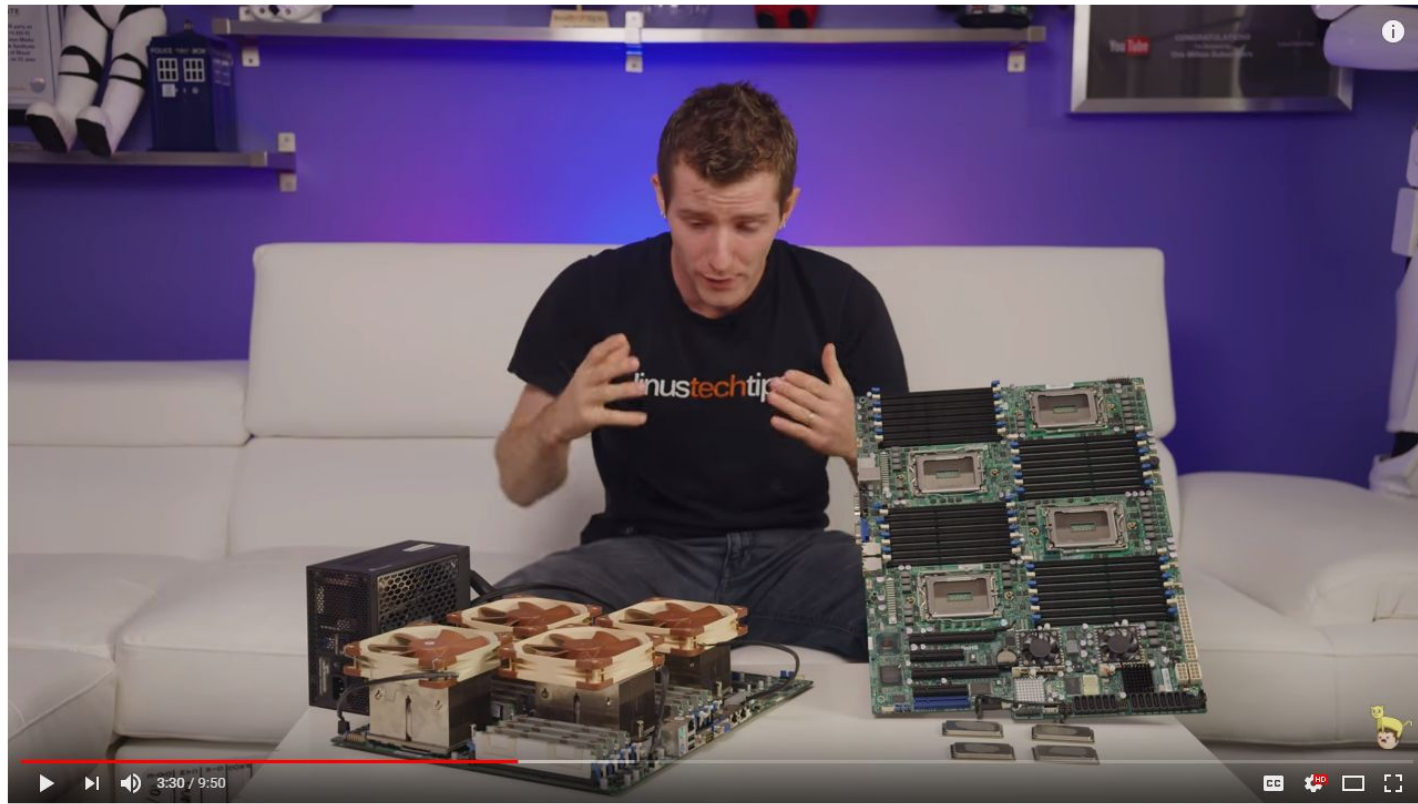
**AMD Threadripper 1950X** (TR4,  
3.40GHz, Unlocked)

★★★★★ 4

## Prozessor



1'839.-

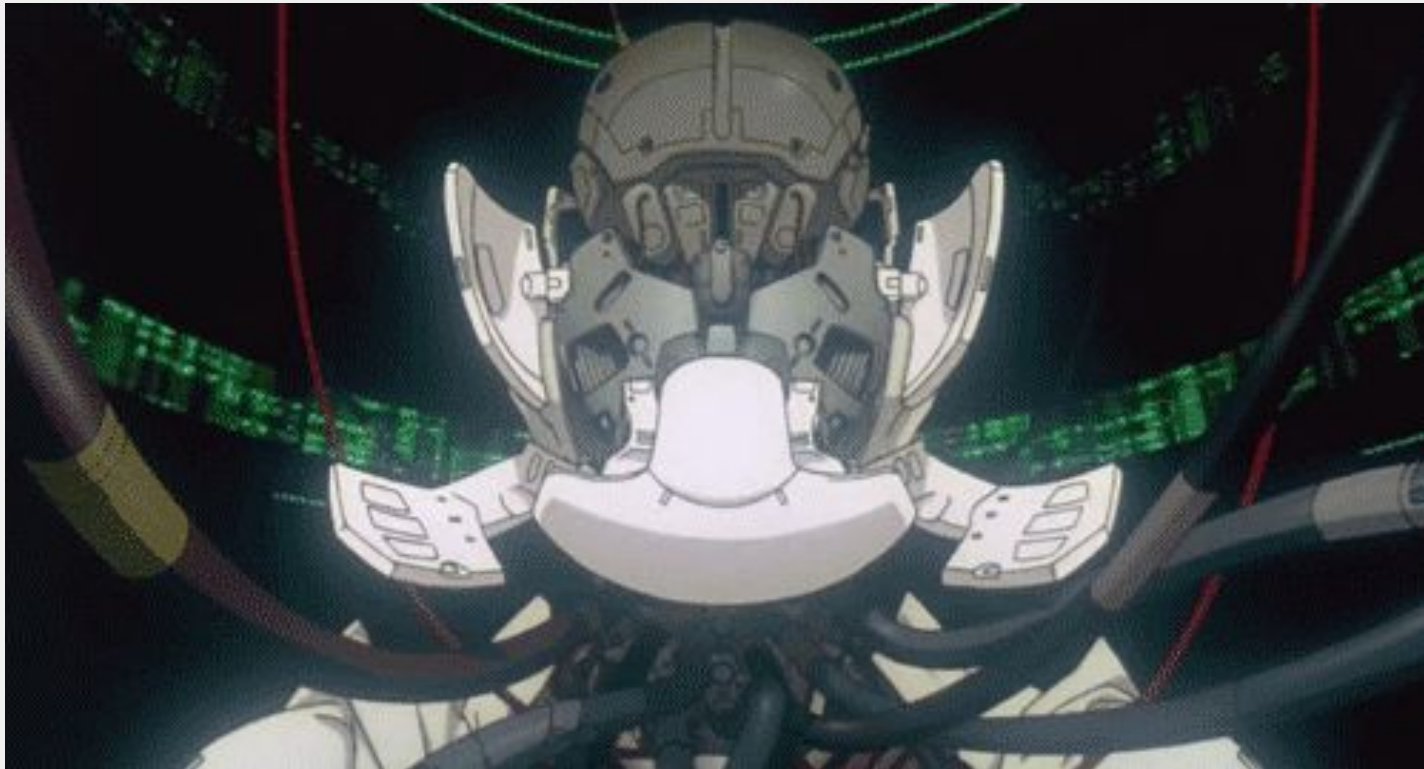
**Intel Core i9-7960X** (LGA 2066,  
2.80GHz, Unlocked)




# Fuzzing Hardware

**What I really need //**




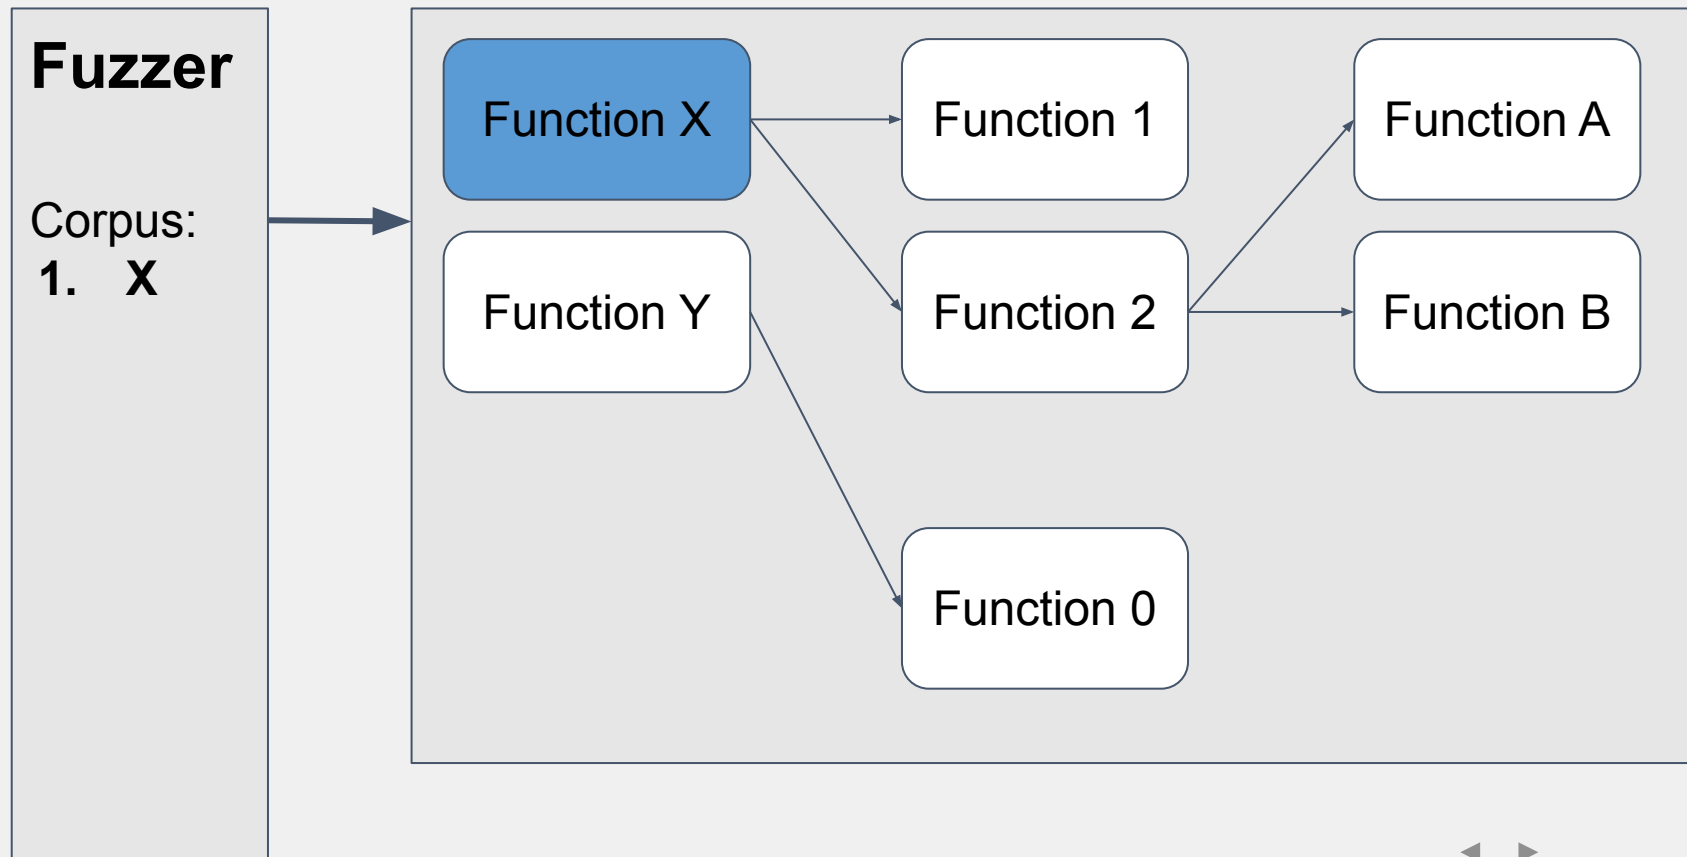


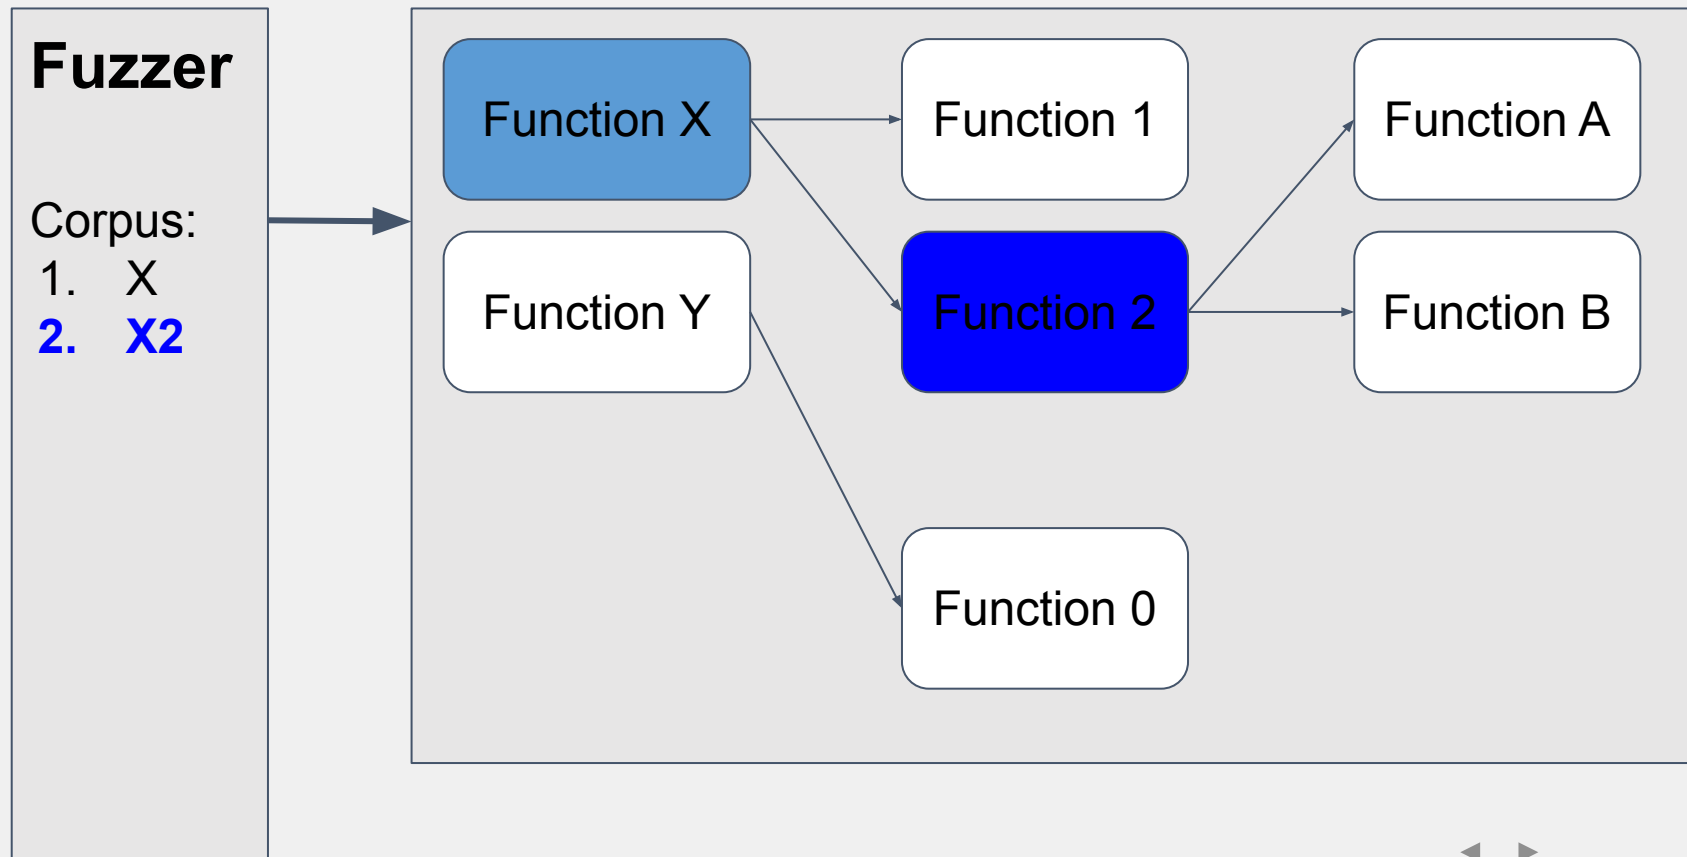
Puppet Master - Sentient Hacking AI, created by Intelligence Unit Section 6 (Ghost in the Shell, 1995)

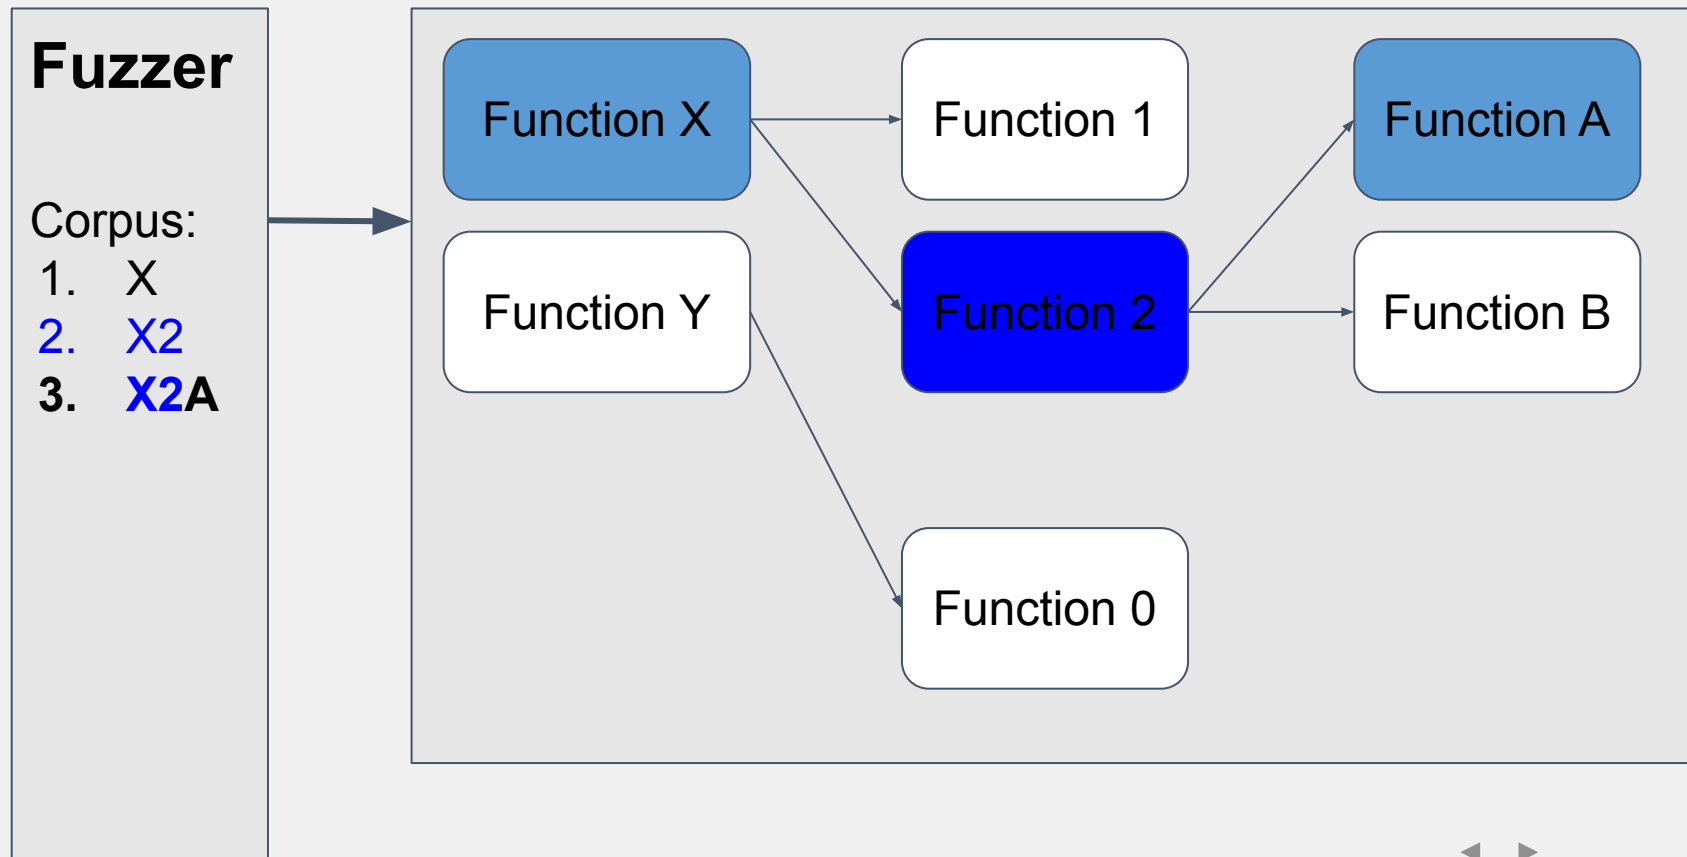


# Feedback Driven Fuzzing





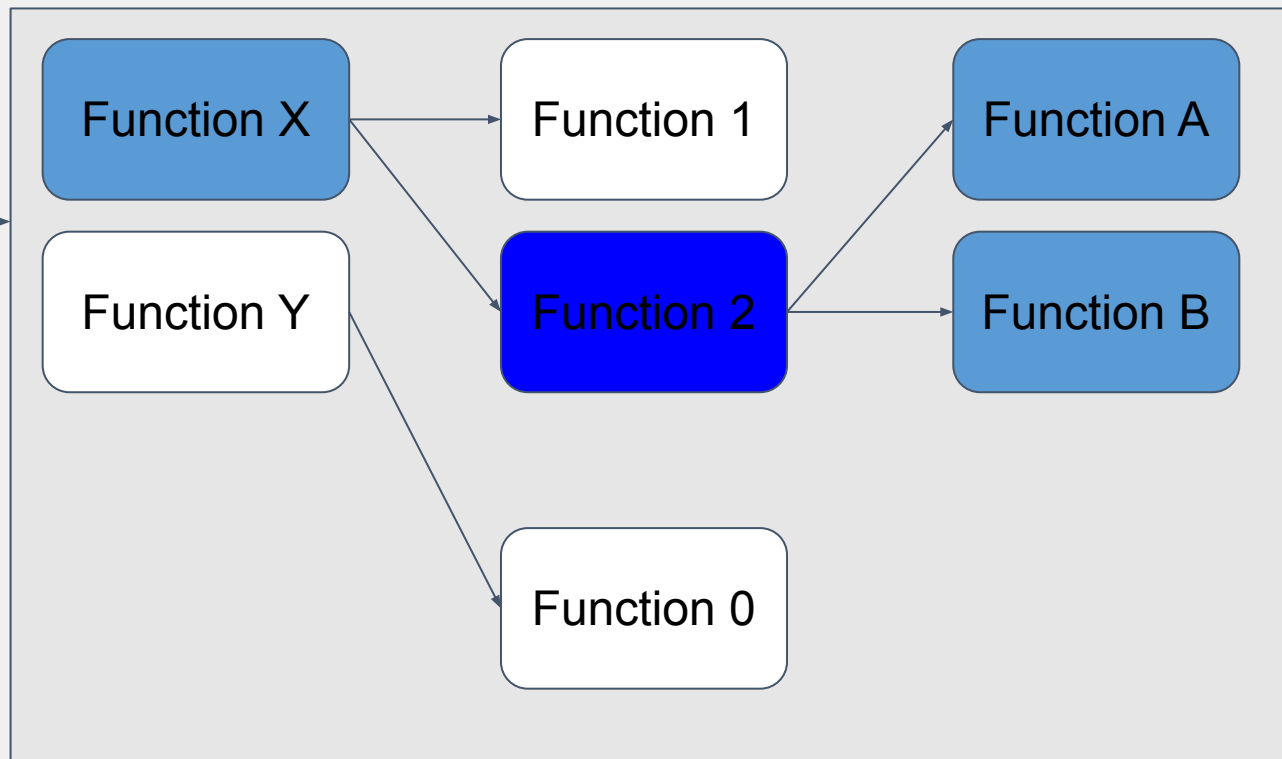




# Fuzzer

Corpus:

1. X
2. X2
3. X2A
4. X2B



```
w = 0;  
x = x + y;  
y = 0;  
if( x > z)  
{  
    y = x;  
    x++;  
}  
else  
{  
    y = z;  
    z++;  
}  
w = x + z;
```

Source Code

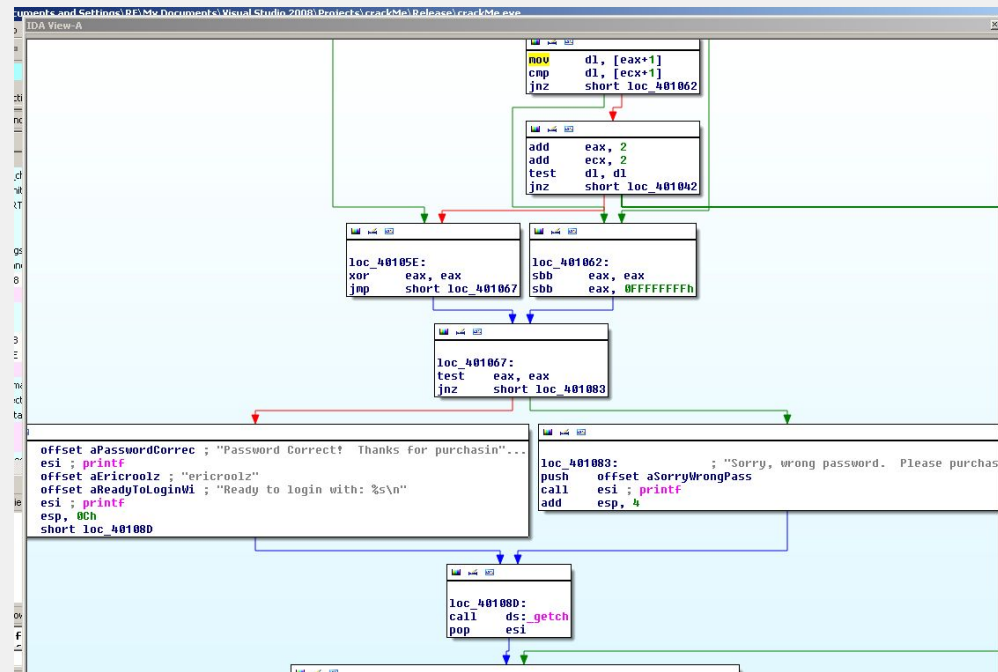
```
w = 0;  
x = x + y;  
y = 0;  
if( x > z)
```

```
y = x;  
x++;
```

```
y = z;  
z++;
```

```
w = x + z;
```

Basic Blocks





## Coverage-guiding in action

The following code wants "ABCD" input:

```
if input[0] == 'A' {  
    if input[1] == 'B' {  
        if input[2] == 'C' {  
            if input[3] == 'D' {  
                slice[input[4]] = 1 // out-of-bounds here  
            }  
        }  
    }  
}
```

Blind generation needs  $O(2^{8^4}) = O(2^{32})$  tries.

Corpus progression:

```
0. {}  
1. {"A"}  
2. {"A", "AB"}  
3. {"A", "AB", "ABC"}  
4. {"A", "AB", "ABC", "ABCD"}
```

Coverage-guided fuzzer needs  $O(4 * 2^8) = O(2^{10})$  tries.

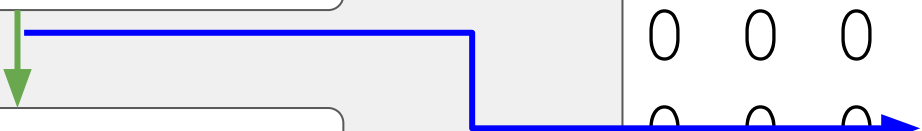
```
if input[0] == 0x41
```

```
if input[1] == 0x42
```

```
if input[2] == 0x43
```

```
if input[3] == 0x44
```

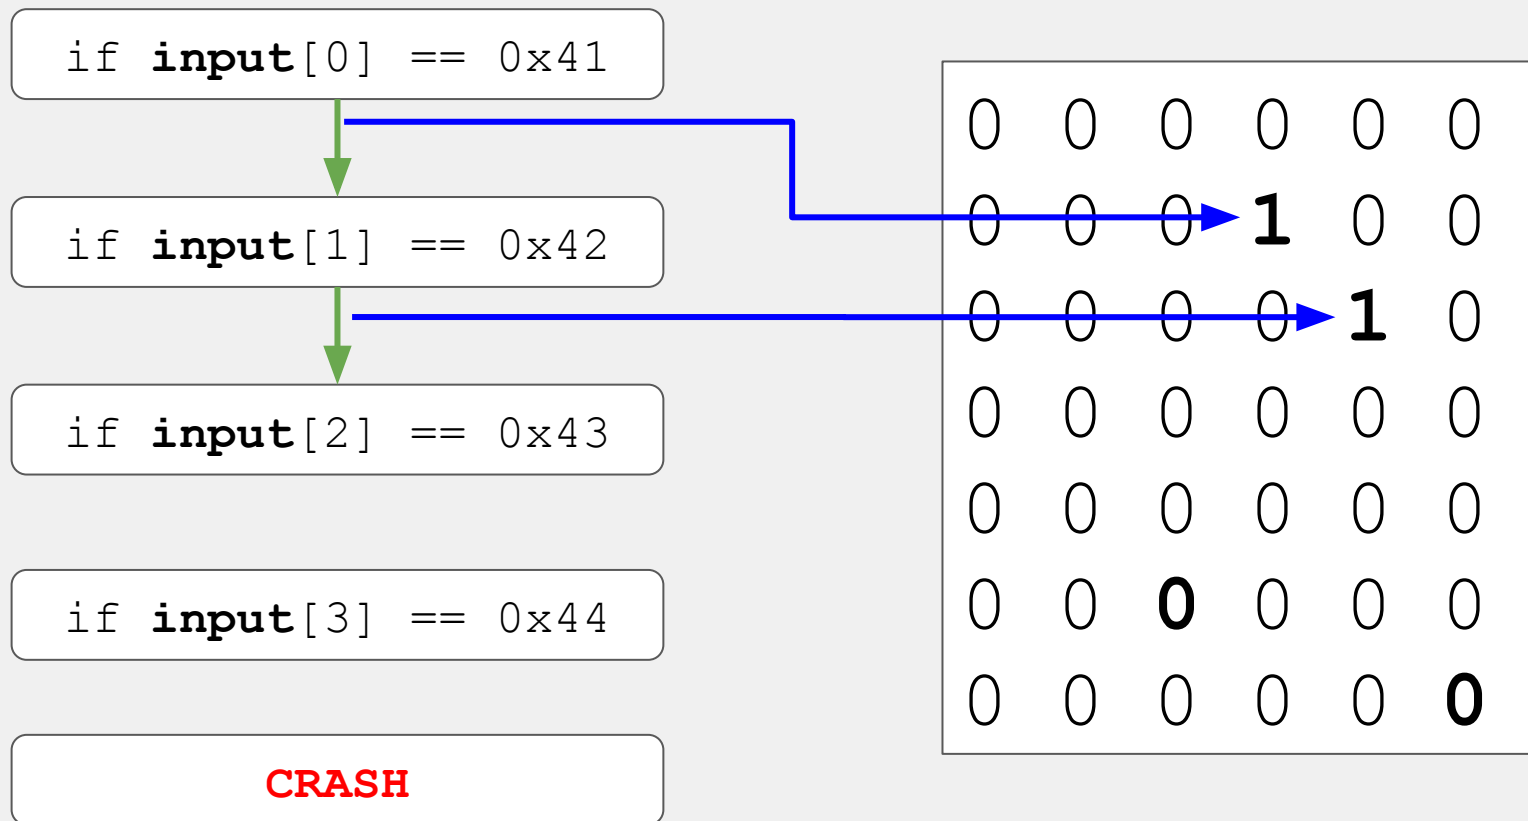
**CRASH**

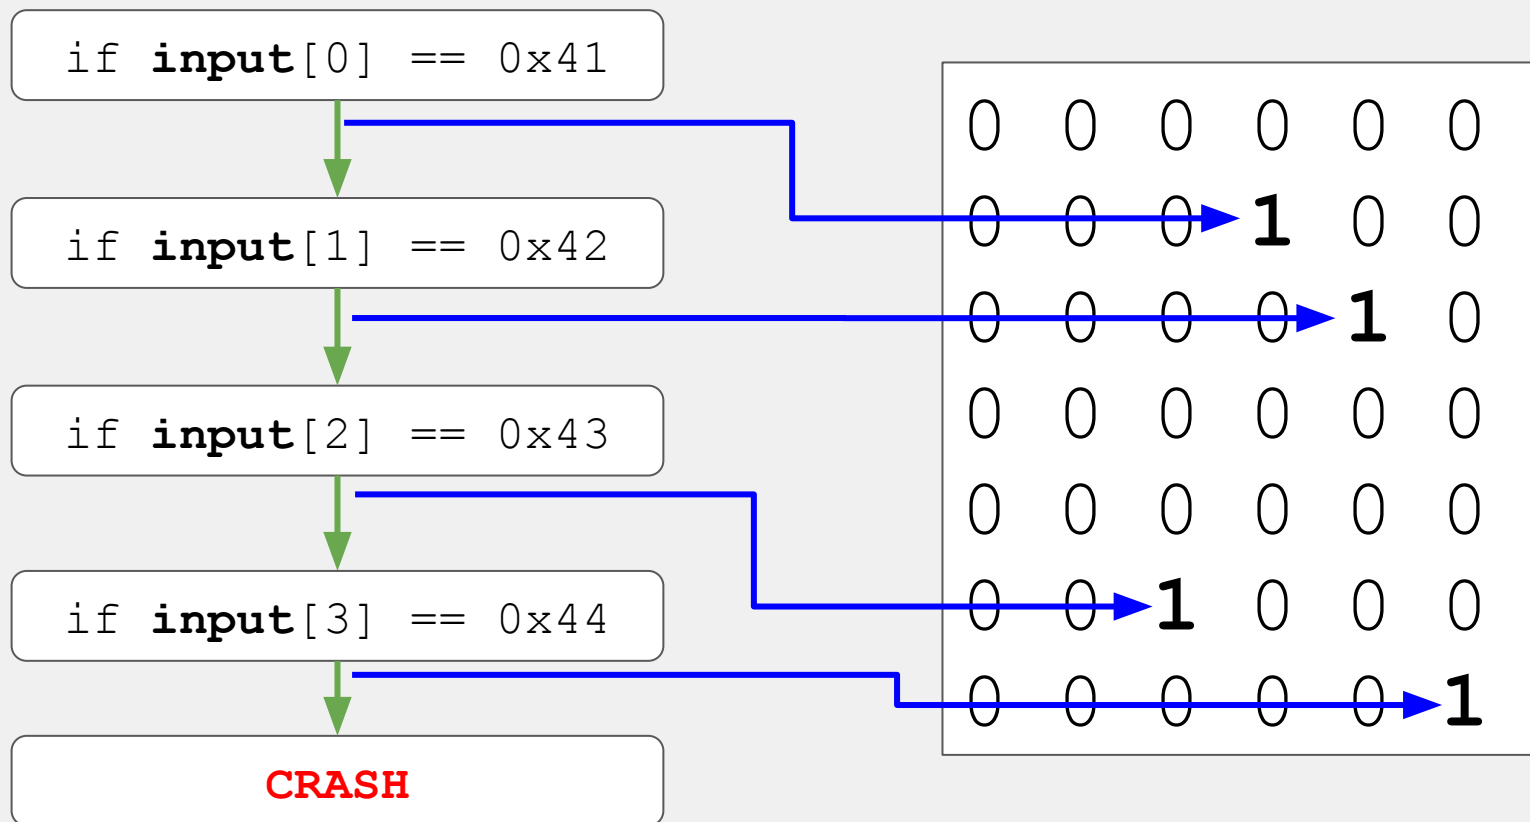


0	0	0	0	0	0
0	0	0	<b>1</b>	0	0
0	0	0	0	<b>0</b>	0
0	0	0	0	0	0
0	0	0	0	0	0
0	0	<b>0</b>	0	0	0
0	0	0	0	0	<b>0</b>

At each basic block, add:

```
cur_location = (block_address >> 4) ^ (block_address << 8);  
  
shared_mem[cur_location ^ prev_location]++;  
  
prev_location = cur_location >> 1;
```





```
109      24 :      } else if (a == 0xFB2Au) { /* SHIN WITH SHIN DOT */
110      0 :          *ab = 0xFB2Cu;
111      0 :          found = true;
112      24 :      } else if (a == 0xFB2Bu) { /* SHIN WITH SIN DOT */
113      0 :          *ab = 0xFB2Du;
114      0 :          found = true;
115      :      }
116      35 :      break;
117      :      case 0x05BFu: /* RAFE */
118      49 :          switch (a) {
119      :              case 0x05D1u: /* BET */
120      12 :                  *ab = 0xFB4Cu;
121      12 :                  found = true;
122      12 :                  break;
123      :              case 0x05DBu: /* KAF */
124      11 :                  *ab = 0xFB4Du;
125      11 :                  found = true;
126      11 :                  break;
127      :              case 0x05E4u: /* PE */
128      14 :                  *ab = 0xFB4Eu;
129      14 :                  found = true;
130      14 :                  break;
131      :          }
132      49 :      break;
133      :      case 0x05C1u: /* SHIN DOT */
134      22 :          if (a == 0x05E9u) { /* SHIN */
135      12 :              *ab = 0xFB2Au;
136      12 :              found = true;
137      10 :          } else if (a == 0xFB49u) { /* SHIN WITH DAGESH */
138      0 :              *ab = 0xFB2Cu;
139      0 :              found = true;
140      :          }
141      22 :      break;
142      :      case 0x05C2u: /* SIN DOT */
143      22 :          if (a == 0x05E9u) { /* SHIN */
144      10 :              *ab = 0xFB2Bu;
145      10 :              found = true;
146      12 :          } else if (a == 0xFB49u) { /* SHIN WITH DAGESH */
147      0 :              *ab = 0xFB2Du;
148      0 :              found = true;
149      :          }
```

Picture Source:

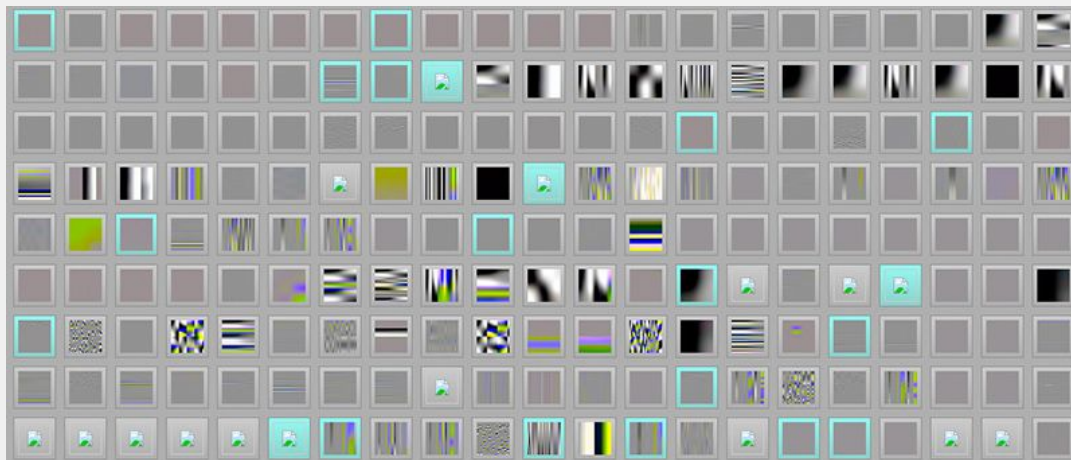
"Circumventing Fuzzing Roadblocks with Compiler Transformations", lafintel

November 07, 2014

## Pulling JPEGs out of thin air

This is an interesting demonstration of the capabilities of [afl](#); I was actually pretty surprised that it worked!

```
$ mkdir in_dir  
$ echo 'hello' >in_dir/hello  
$ ./afl-fuzz -i in_dir -o out_dir ./jpeg-9a/djpeg
```



# Feedback-driven fuzzing

Honggfuzz is capable of performing feedback-guided (code coverage driven) fuzzing. It can utilize the following sources of data:

- (Linux) Hardware-based counters (instructions, branches)
- (Linux) Intel BTS code coverage (kernel  $\geq 4.2$ )
- (Linux) Intel PT code coverage (kernel  $\geq 4.2$ )
- Sanitizer-coverage instrumentation ( `-fsanitize-coverage=bb` )
- Compile-time instrumentation ( `-finstrument-functions` or `-fsanitize-coverage=trace-pc[-guard],indirect-calls,trace-cmp` or both)

Developers should provide the initial file corpus which will be gradually improved upon. It can even comprise of a single 1-byte initial file, and honggfuzz will try to generate better inputs starting from there.



google / honggfuzz

Watch 93 Star 880 Fork 197

<> Code Issues 2 Pull requests 0 Projects 0 Insights

Branch: master honggfuzz / docs / FeedbackDrivenFuzzing.md Find file Copy path

robertswiecki cmdline: make the compile-time instrumentation the default option, ca... 930e12f 11 days ago

3 contributors

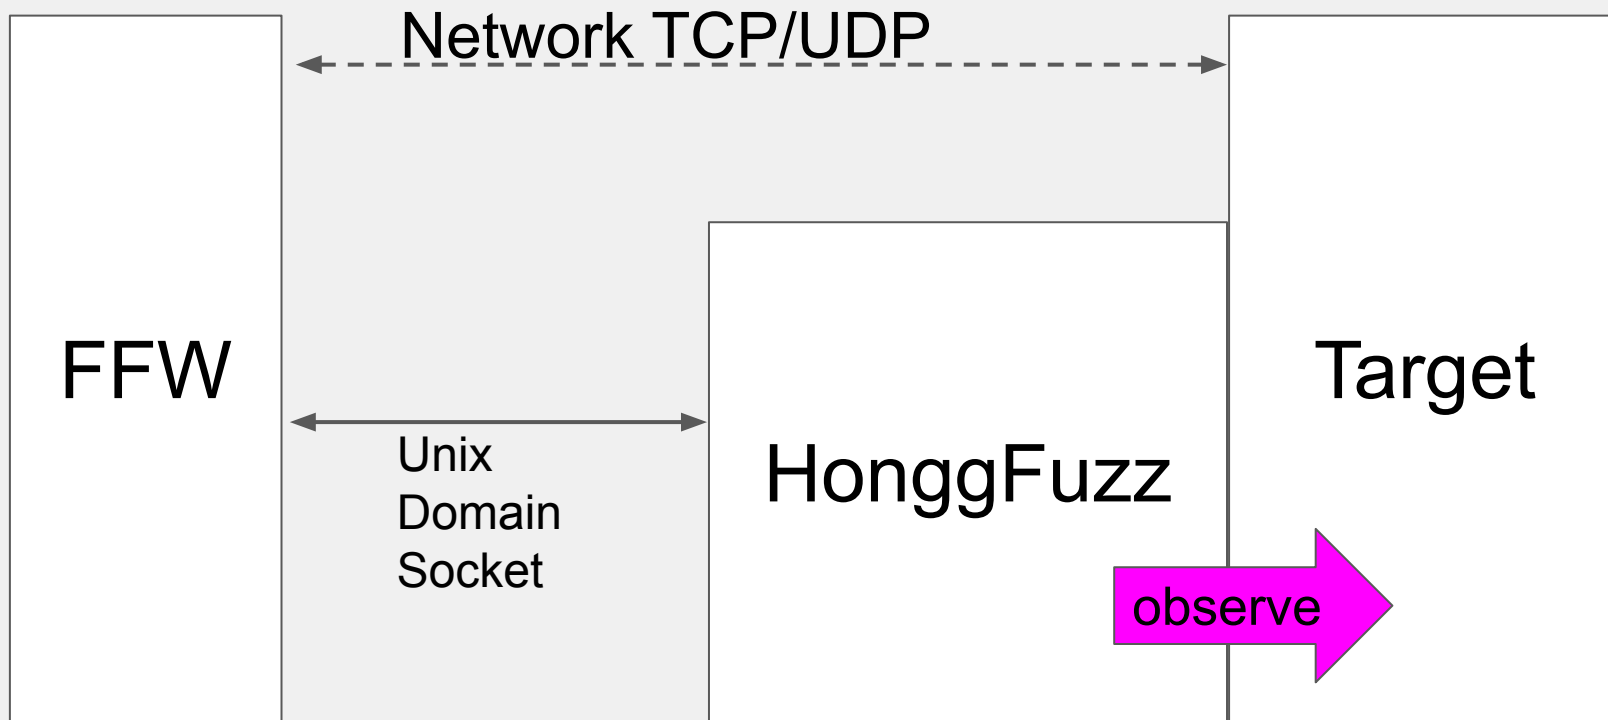
244 lines (204 sloc) 11.6 KB Raw Blame History

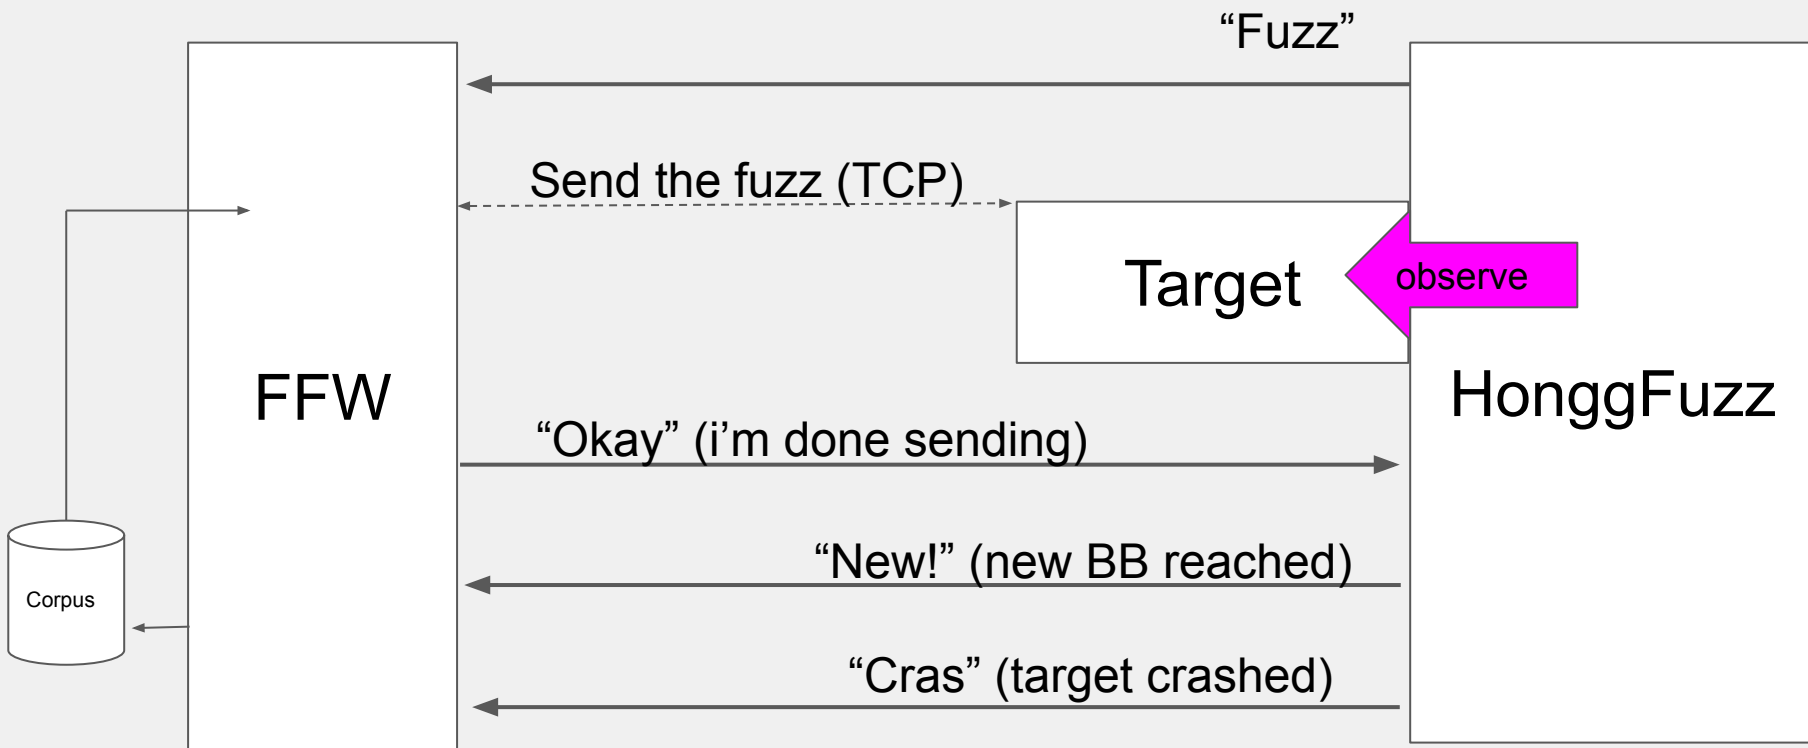
## Feedback-driven fuzzing

Honggfuzz is capable of performing feedback-guided (code coverage driven) fuzzing. It can utilize the following sources of data:

- (Linux) Hardware-based counters (instructions, branches)
- (Linux) Intel BTS code coverage (kernel >= 4.2)
- (Linux) Intel PT code coverage (kernel >= 4.2)
- Sanitizer-coverage instrumentation ( `-fsanitize-coverage=bb` )
- Compile-time instrumentation ( `-finstrument-functions` or `-fsanitize-coverage=trace-pc[-guard],indirect-calls,trace-cmp` or both)

Developers should provide the initial file corpus which will be gradually improved upon. It can even comprise of a single 1-byte initial file, and honggfuzz will try to generate better inputs starting from there.





FFW  
Mongoose 6.9  
Feedback driven fuzzing with honggfuzz  
Demo

Praise  
the  
Demo  
Gods

┌  
Performance  
└

About **30 iterations/s** per thread :-)

AFL: Threads = core\_count / 2

FFW: Threads = core\_count \* 2

- 30 iterations/s \* 32 \* 2 = ~ 2000 (okish)

## Parallel fuzzing -> parallel processes

```
-listen %(port)
```

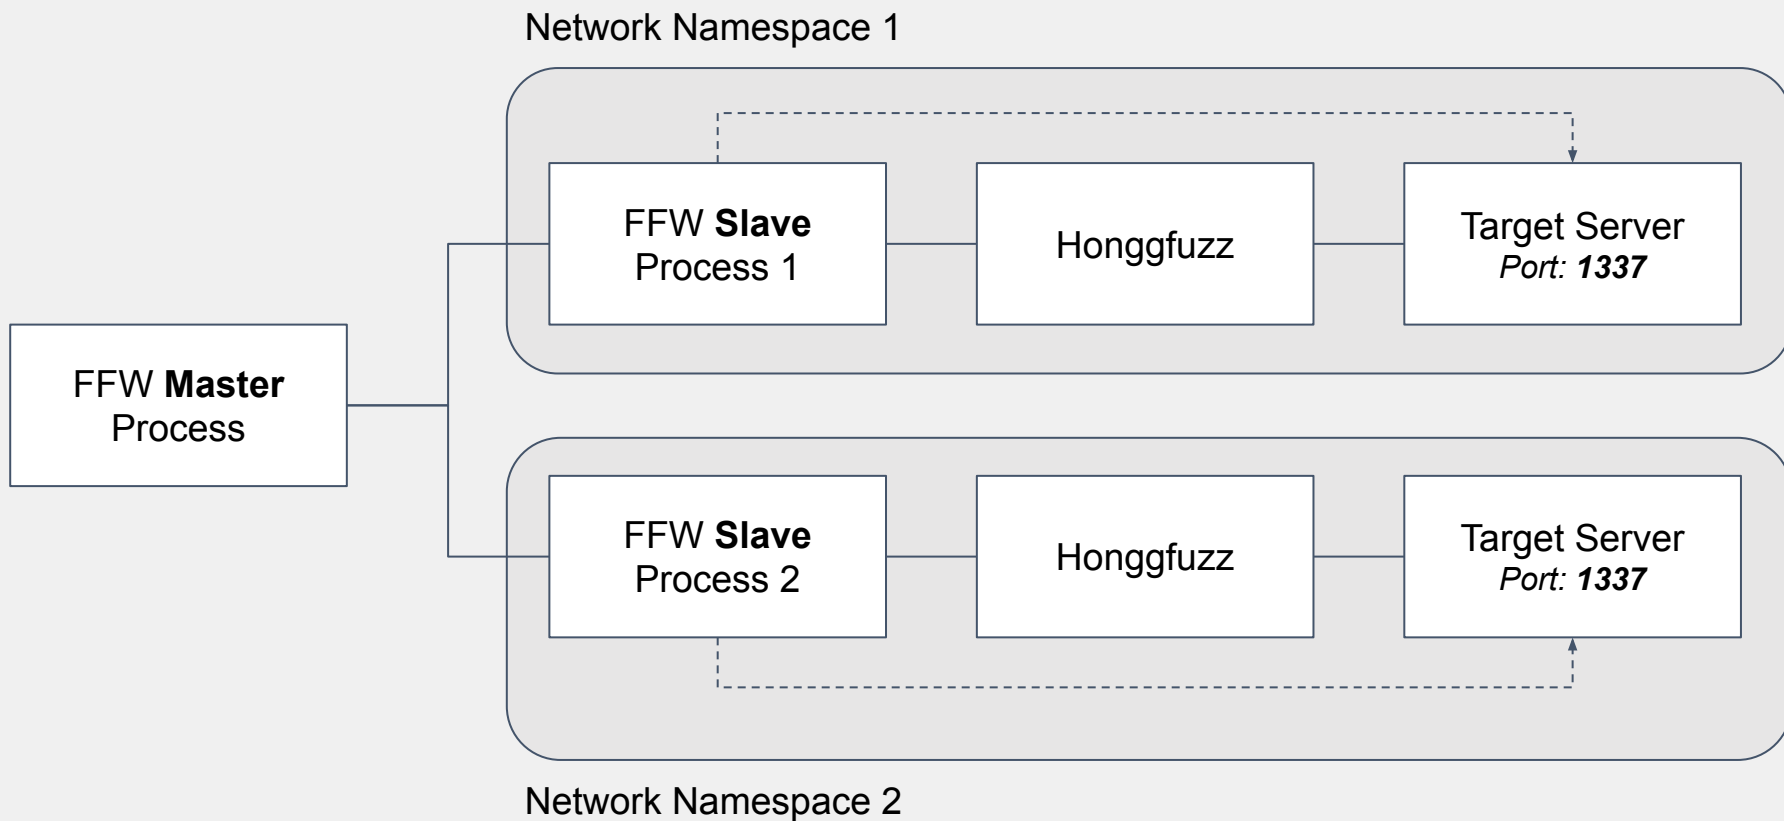
Port directly  
as Argument

```
-config srvcfg-%(port).conf
```

Port via  
Config File

```
usenamespace: "True"
```

Namespaces!  
yay!







# Fuzzing Results



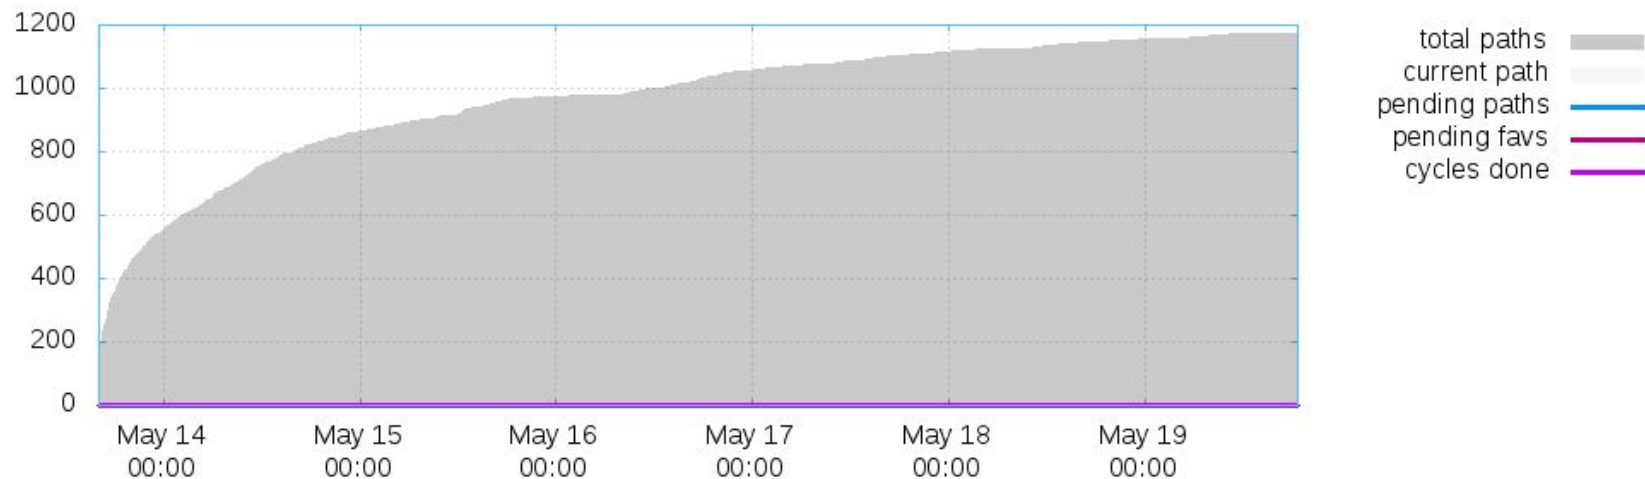
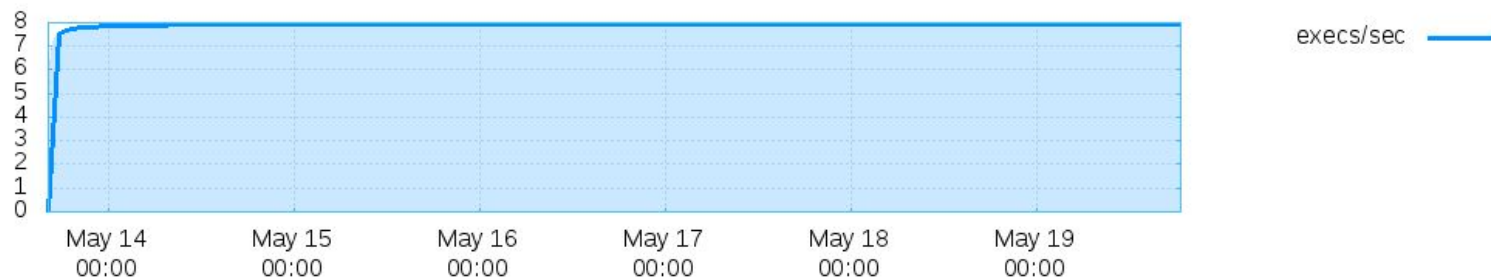
## Bugs:

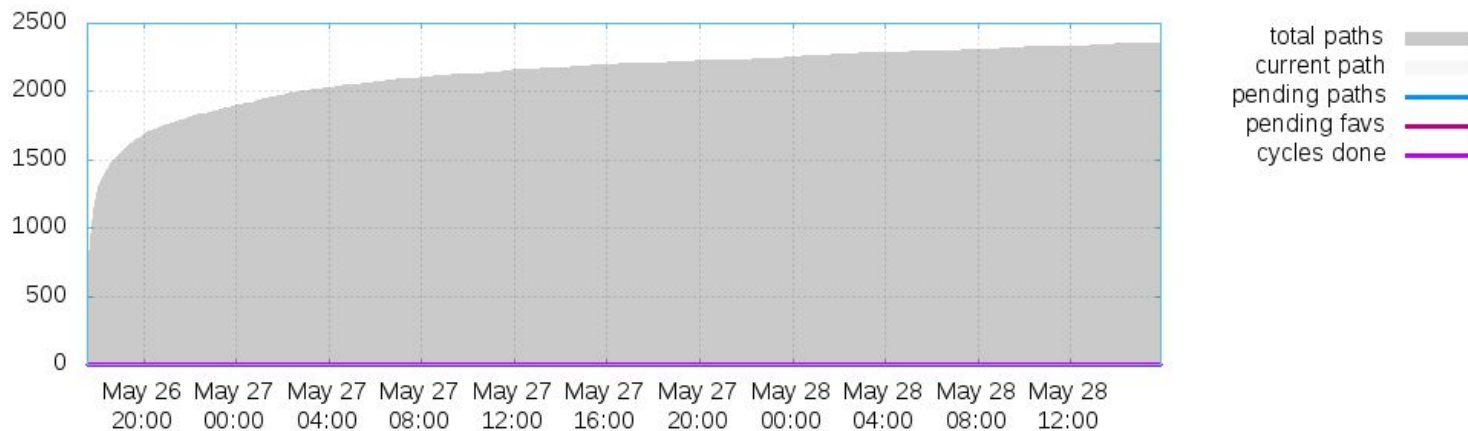
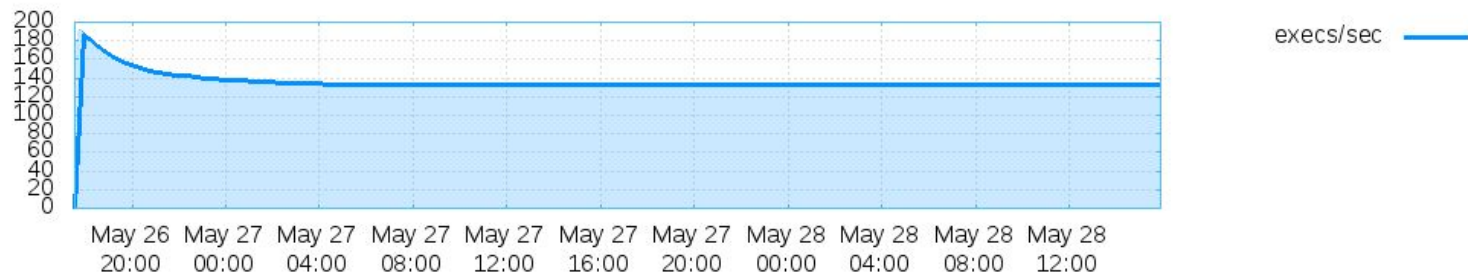
- Mongoose MQTT 6.8
- Mongoose MQTT 6.9
- libcoap
- libiec61850
- vnc4server
- libmodbus

Check docker dobin/ffw:0.2

## No Bugs:

- *Mongoose Web*
- *Mongoose DNS*
- *Mongoose DHCP*
- Synergy
- Ngircd
- Inspircd
- unrealircd
- portmapper





# Conclusion



01

**INTERCEPT**

02

**RETRANSMIT (fuzzed)**

03

**DETECT**

**Crashes**

**New Basic Blocks**



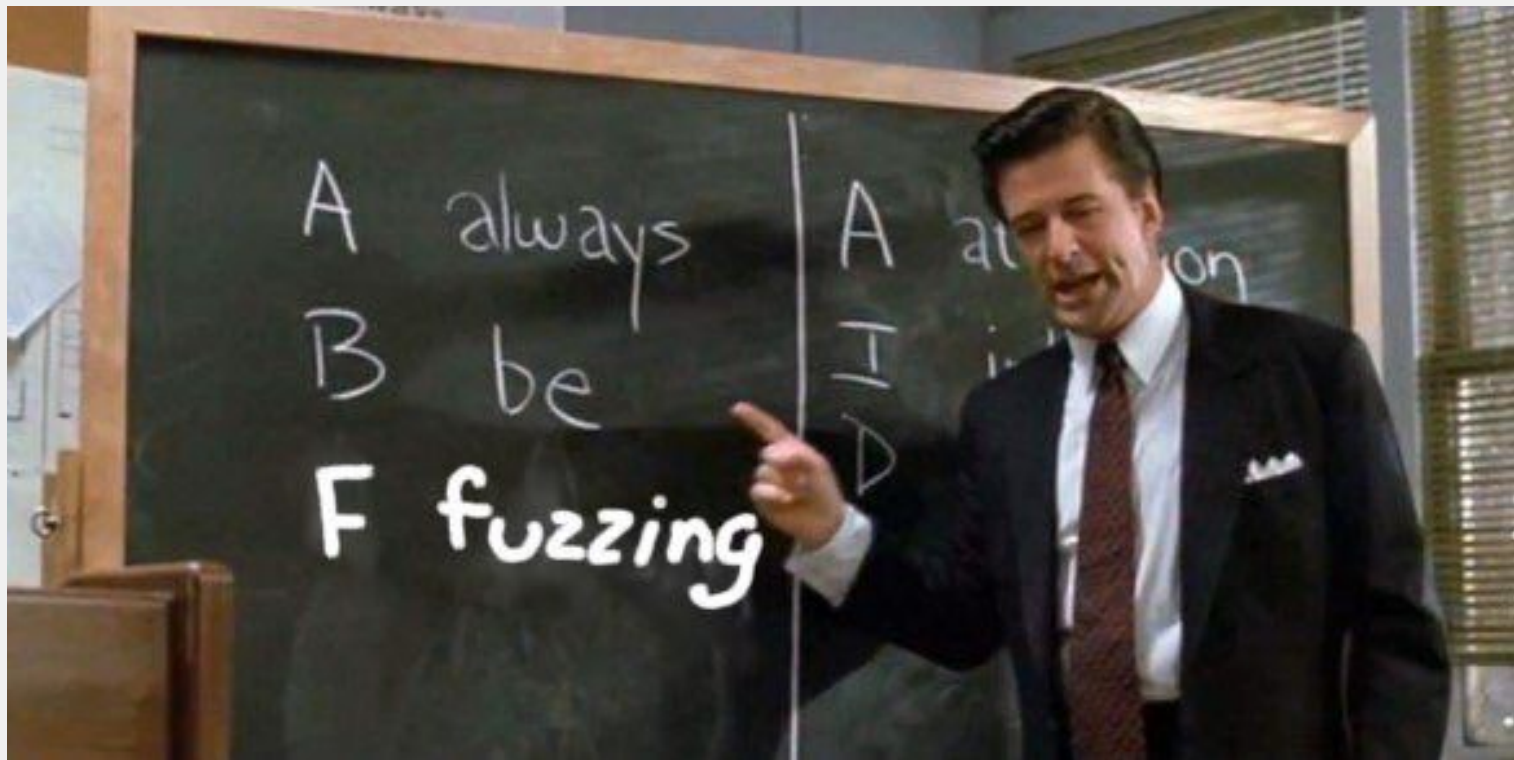
Swat low hanging bugs



Use time before *real* fuzzing



Quick-Check of a server





Code: **[github.com/dobin/ffw](https://github.com/dobin/ffw)**

Me: **[twitter.com/dobinrutis](https://twitter.com/dobinrutis)**

Blog: **[dobin.ch/blog/tag/ffw](https://dobin.ch/blog/tag/ffw)**

0days: **[github.com/dobin/ffw-docker](https://github.com/dobin/ffw-docker)**

Reddit: **[reddit.com/r/fuzzing](https://reddit.com/r/fuzzing)**

### **Get started:**

```
docker run -ti --privileged -lxc-conf="aa_profile=unconfined" dobini/ffw:0.2
```

# Q&A

### Not in this presentation:

- Taint Tracking / Analysis
- SMT Solvers
- Automated root cause identification
- Automated exploit generation
- Automated patch generation
- Machine Learning / AI
- Kernel Fuzzing
- Syscall Fuzzing
- Instruction-Set Fuzzer
- Structure-aware fuzzing
- Anti-Fuzzing
- Binary Instrumentation (Intel PIN, DynamoRio)
- Blockchain (Smart Contracts / Ethereum)
- Symbolic Execution / Concolic Execution (Angr, Manticore)
- Compiler Transformations / Deoptimization



## Fuzzotron

---

Available via <https://github.com/denandz/fuzzotron>. "Fuzzotron is a simple network fuzzer supporting TCP, UDP and multithreading."

Support network fuzzing, also uses Radamsa. Can use coverage data, but it is experimental.

Con's:

- Does not restart target server
- Unreliable crash detection
- Experimental code coverage

## Mutiny

---

Available via <https://github.com/Cisco-Talos/mutiny-fuzzer>. "The Mutiny Fuzzing Framework is a network fuzzer that operates by replaying PCAPs through a mutational fuzzer."

Con's:

- No code coverage
- Only one commit (no development?)
- Rudimentary crash detection

Future

**Fitness:**

Fuzz which damn input?

- [] Input Corpus
  - [] Corpus Network Message

01 # Crashes

02 # Hangs

03 # Children

04 Latency

## Generic Mutation Algorithm

Per InputCorpus



### Crossover

(copy partially to other network messages)



### Mutation

(re-order network messages)



### Swap

(swap messages)

**Crash Analysis**

**Code Coverage in IDE**

**Adjust network timeout  
dynamically**

**Client program fuzzing**



I wrote a vulnerability scanner that abstracts all the predicates in a binary, traverses the callgraph and generates phormulaes to run then with a SMT solver. I found 1 vuln in 3 days with this tool.



He wrote a dumb ass fuzzer and found 5 vulns in 1 day.

Good thing I'm not a n00b like that guy.



The reason for all this:

C  
C++

# Feedback Mechanism

	Closed-Source	Kernel	Stable	Fast
Compile-Time Instrumentation	✗	✓	✓	+++
Static Rewriting	✓	-	✗	+++
Dynamic Binary Instrumentation	✓	_*	✓	-
Emulation	✓	✓	✓	--
Intel Branch Trace Store	✓	✓	✓	+
Intel Processor Trace	✓	✓	✓	++++

\* Peter Feiner, et al., DRK: DynamoRIO as a Linux Kernel Module